

Chapter 1 Classification

1.1 Classification and Representation

1.1.1 Classification

Classification is to predict the variable y into discrete values, in other words, an assignment to classify features into classes. For example, divide emails by spam or not.

We can set a list of threshold of some parameters of hypothesis to divide the dataset into different classes.

But if we apply linear regression, some extreme data will have a nonnegligible effect on the hypothesis.

As we all know, the hypothesis usually gives a continuous value, while the classification problem gives discrete values. But apply hypothesis and map values into discrete values sometimes cannot work well.

Here we will focus on the binary classification which has a positive class and a negative class.

1.1.2 Hypothesis Representation

A logistic regression model wants its $h_{\theta}(x) \in [0, 1]$. We could apply **Sigmoid Function**:

$$g(z) = \frac{1}{1 + e^{-z}}$$

It's like




Figure 1.1: Sigmoid Function

Then, we get:

$$h_{\theta}(x) = g(\theta^T x)$$

Let's interpret the hypothesis' output: $h_{\theta}(x)$ is the estimated probability that $y = 1$ on input x . That is **probability that $y = 1$, given a x parameterized by θ** :

$$h_{\theta}(x) = P(y = 1|x; \theta)$$

1.1.3 Decision Boundary

Why Sigmoid Function makes sense? When $h_{\theta}(x) > 0.5$ we predict $y = 1$ and predict $y = 0$ in else condition. And it means that $\theta^T x > 0$ or not.

$$h_{\theta}(x) = g(\theta^T x)$$

By Sigmoid Function we can classify the hypothesis by it's values.

When we get 2 variables, like $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$ and we still make binary classification, wo we can still predict $y = 1$ if $h_{\theta}(x) \geq 0$ and $y = 0$ if $h_{\theta}(x) \leq 0$.

Further more, how about we need a non-linear decision boundaries, like a circle or a a ellipse? Just use non-linear functions! For a ellipse:

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

And we define the boundary like:

$$\begin{cases} y = 1, & \text{if } x_1^2 + x_2^2 \geq 1 \\ y = 0, & \text{else} \end{cases}$$

1.2 Logistic Regression Model

1.2.1 Cost Function

For a classification problem, we need the cost function, too. This is related to how to choose the parameters.

In chapters before, we defined as a sum of squared error:

$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

For a simple cost function, it's non-convex so it does not guarantee to converge.

In logistic regression cost function, we define:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

If the hypothesis gives an approximate output as y , the cost is rather small; but when it goes to the other end, the cost will grow in a very fast speed.

And this cost function could be convex in our problem.

1.2.2 Simplified Cost Function and Gradient Descent

Let's re-write the cost function in a more compact way as:

$$\text{Cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x), y) - (1 - y) \log(1 - h_{\theta}(x))$$

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[y^{(i)} \log(h_{\theta}(x^{(i)}), y^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

To fit parameters θ , we need to minimize the $J(\theta)$. Then we can apply the gradient descent as in previous chapters.

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=0}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}, \text{ where } \theta = 0, 1, 2, \dots, n$$

Note: the $h_{\theta}(x)$ is different! Now we have $h_{\theta}(x) = 1/(1 + e^{-\theta x})$.

1.2.3 Advanced Optimization

With this section, we can get logistic regression run more quickly. Like, gradient descent, conjugate gradient, BFGS¹, L-BFGS²

We can just simply call `fminunc` to get gradient descent.

1.3 Multi-class Classification

1.4 Solve Overfitting

¹Broyden-Fletcher-Goldfarb-Shanno algorithm,

²Limited-memory BFGS